



Data Sampling, Visualization, Learning and Classification

Machine Learning – Laboratory

Battista Biggio

Department of Electrical and Electronic Engineering
University of Cagliari, Italy

Exercise 1

- Consider the function `make_gaussian_dataset` defined in the previous lab session
- Extend it to handle:
 1. more than two dimensions
 2. more than two classes
 3. non-isotropic Gaussians
 - covariance matrix is a positive-definite matrix, not necessarily proportional to the identity matrix (namely, features are correlated!)

Exercise 1 – This is the starting point...

```
import numpy as np

def make_gaussian_dataset(n0, n1, mu0, mu1):
    """ Creates a 2-class 2-dimensional Gaussian dataset. """
    d = 2 # hard-coded for convenience, we will improve this later on

    x0 = np.random.randn(n0, d) + mu0 # uses broadcasting...
    x1 = np.random.randn(n1, d) + mu1

    # sample labels
    y0 = np.zeros(n0)
    y1 = np.ones(n1)

    # concatenate data and labels
    x = np.vstack((x0, x1))
    y = np.hstack((y0, y1))

    return x, y

# generate data with 10 samples/class, and means [-1,-1], [1, 1]
xn, yn = make_gaussian_dataset(10, 10, [-1, -1], [+1, +1])
print('xn: ', xn)
print('yn: ', yn)
```

Exercise 1: Solution

```
def make_gaussian_dataset(n, mu):  
    """  
    Creates a k-class d-dimensional Gaussian dataset.  
    :param n: vector containing the number of samples for each class  
    :param mu: matrix containing the mean vector for each class  
    :return: x,y, the gaussian dataset  
    """  
  
    n = np.array(n)  # convert to np.array if list is passed as input  
    mu = np.array(mu)  
    n_classes = mu.shape[0]  # number of classes  
    n_features = mu.shape[1]  # number of features  
    n_samples = n.sum()  # total number of samples  
  
    x = np.zeros(shape=(n_samples, n_features))  
    y = np.zeros(shape=(n_samples,))  
  
    start_index = 0  
    for i in xrange(n_classes):  
        x_tmp = np.random.randn(n[i], n_features) + mu[i, :]  # broadcasting...  
        x[start_index:start_index + n[i], :] = x_tmp  
        y[start_index:start_index + n[i]] = i  
        start_index += n[i]  
    return x, y
```

Exercise 1: Solution

- This is still not considering different covariance matrices per class

```
import numpy as np

def make_gaussian_dataset(n, mu):
    [...]

    # generate data
    xn, yn = make_gaussian_dataset([10, 5, 2], [[-1, -1],
                                                  [+1, -1],
                                                  [-1, +1]])
```

- How to extend it to use a different covariance matrix per class?
`make_gaussian_dataset(n, mu, cov)`?

Exercise 2

Define a function that plots a dataset using a different color for each class:

```
plot_dataset (x, y, feat_1=0, feat_2=1)
```

Hints:

```
import matplotlib.pyplot as plt
```

```
plt.scatter(x1, x2, color='r')
```

plots the point (x1,x2) as a red point

Colors are: ['k', 'b', 'r', 'g', 'c', 'm', 'y']

```
bool_class0=(y==0) # select samples belonging to class 0
```

Other useful functions:

```
plt.xlabel(), plt.ylabel(), plt.legend(), plt.show()
```

Exercise 2: Solution

```
import matplotlib.pyplot as plt

def plot_dataset(x, y, feat_1=0, feat_2=1):
    n_classes = len(np.unique(y))
    colors = ['r', 'b', 'k', 'g', 'c', 'm', 'y']

    for y0 in xrange(n_classes):
        x0 = x[y == y0, feat_1]  # y0 is the current class in the loop
        x1 = x[y == y0, feat_2]
        plt.scatter(x0, x1, c=colors[y0], label='class ' + str(y0))


    plt.legend()
    plt.xlabel('feature x' + str(feat_1))
    plt.ylabel('feature x' + str(feat_2))

    return
```

Exercise 2: Solution

generate data

```
xn, yn = make_gaussian_dataset([50, 50, 50], [[-5, -5], [+5, -5], [-5, +5]])  
plot_dataset(xn, yn, 0, 1)
```



Learning and Classification

What's next? Learning and Classification

- Now we can sample data and visualize it in two dimensions
- The goal of the next exercises is to implement a simple classifier
 - The Nearest Mean Classifier (NMC)
- We will implement its learning and classification procedures

Ex. 3: NMC – Learning & Classification

- During the learning phase, NMC is given a training set consisting of pairs (x, y) of samples along with their labels
- For each class y_0 in y
 - NMC estimates the mean of the samples in class y_0
 - stores the mean vector (centroid)
- During classification, NMC assigns the current test sample x to the class whose mean vector (centroid) is the closest one to x
- Implement the functions
 - `centroids = fit(x, y)`, corresponding to the learning phase, and
 - `y_pred, distances = predict(x, centroids)`, corresponding to the classification phase, where `y_pred` is the label of the predicted class, and `distances` are the distance values w.r.t the centroids

Exercise 3: Solution

```
import numpy as np

def fit(x, y):
    n_classes = np.unique(y).size
    n_features = x.shape[1]

    centroids = np.zeros(shape=(n_classes, n_features))
    for k in xrange(n_classes):
        centroids[k] = x[y == k, :].mean(axis=0)
    return centroids

def predict(x, centroids):
    n_samples = x.shape[0]
    n_classes = centroids.shape[0]
    distances = np.zeros(shape=(n_samples, n_classes))

    for k in xrange(n_classes):
        distances[:,k] = np.linalg.norm(x-centroids[k, :], axis=1)
    y_pred = np.argmin(distances, axis=1)

    return y_pred, distances
```

Let's create a class

```
class CNearestMeanClassifier(object):
    """Class implementing a nearest mean classifier"""

    def __init__(self):
        self._centroids = None
        return

    def fit(self, x, y):
        n_classes = np.unique(y).size
        n_features = x.shape[1]
        centroids = np.zeros(shape=(n_classes, n_features))
        for k in xrange(n_classes):
            centroids[k] = x[y == k, :].mean(axis=0)
        self._centroids = centroids
        return

    def predict(self, x):
        n_samples = x.shape[0]
        n_classes = self._centroids.shape[0]
        distances = np.zeros(shape=(n_samples, n_classes))
        for k in xrange(n_classes):
            distances[:, k] = np.linalg.norm(x - self._centroids[k, :], axis=1)
        y_pred = np.argmin(distances, axis=1)
        return y_pred, distances
```

Class Properties

- Python decorator to access class private members
 - See also 'setters'

```
class CNearestMeanClassifier(object):  
    """Class implementing a nearest mean classifier"""  
  
    def __init__(self):  
        self._centroids = None  
        return  
  
    @property  
    def centroids(self):  
        return self._centroids  
  
    [...]
```

Scikit-learn Classifiers

- Check http://scikit-learn.org/stable/supervised_learning.html
- NearestCentroid implements our CNearestMeanClassifier
 - <http://scikit-learn.org/stable/modules/neighbors.html#nearest-centroid-classifier>

The screenshot shows the Scikit-learn website documentation. At the top, there's a navigation bar with links for Home, Installation, Documentation, and Examples. A search bar is also present. Below the navigation bar, there's a section for "1. Supervised learning" and a sub-section for "1.1. Generalized Linear Models". The sub-section lists various models: 1.1.1. Ordinary Least Squares, 1.1.2. Ridge Regression, 1.1.3. Lasso, 1.1.4. Multi-task Lasso, 1.1.5. Elastic Net, 1.1.6. Multi-task Elastic Net, 1.1.7. Least Angle Regression, and 1.1.8. LARS Lasso. On the left side, there's a sidebar with links for Previous, Next, and Up, and a note about the documentation version (0.18.1). A diagonal banner on the right says "Fork me on GitHub".

scikit-learn

Home Installation Documentation Examples

Google Custom Search Search x

powered by Google

Previous User Guide Next 1.1. General... User Guide Up User Guide

This documentation is for scikit-learn version 0.18.1 — Other versions

If you use the software, please consider citing scikit-learn.

1. Supervised learning

1.1. Generalized Linear Models

- 1.1.1. Ordinary Least Squares
 - 1.1.1.1. Ordinary Least Squares Complexity
- 1.1.2. Ridge Regression
 - 1.1.2.1. Ridge Complexity
 - 1.1.2.2. Setting the regularization parameter: generalized Cross-Validation
- 1.1.3. Lasso
 - 1.1.3.1. Setting regularization parameter
 - 1.1.3.1.1. Using cross-validation
 - 1.1.3.1.2. Information-criteria based model selection
- 1.1.4. Multi-task Lasso
- 1.1.5. Elastic Net
- 1.1.6. Multi-task Elastic Net
- 1.1.7. Least Angle Regression
- 1.1.8. LARS Lasso

Fork me on GitHub

Ex. 4: Visualizing the decision regions

```
def plot_decision_regions(x, y, classifier, resolution=0.02):
    # setup marker generator and color map
    colors = ('red', 'blue', 'lightgreen', 'black', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = x[:, 0].min() - 1, x[:, 0].max() + 1
    x2_min, x2_max = x[:, 1].min() - 1, x[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z, score = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    plt.contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())

    # plot class samples
    plot_dataset(x, y)
    return
```


Exercise 4: Solution

```
from src.prilib import CNearestMeanClassifier, \
    make_gaussian_dataset, plot_decision_regions
import matplotlib.pyplot as plt


x, y = make_gaussian_dataset([50, 50, 50], [[-5, -5],
                                             [+5, -5],
                                             [-5, +5]])

classifier = CNearestMeanClassifier()
classifier.fit(x, y)

plot_decision_regions(x, y, classifier)

# plot centroids
plt.scatter(classifier.centroids[:, 0],
            classifier.centroids[:, 1],
            marker='x', color='k')

plt.show()
```



Ex. 5: Testing performance on unseen data

- To assess classifier performance, one should estimate the classification error on never-before-seen data
 - The training data should not be used to this end, as it provides an optimistic estimate of the real performance!
- Therefore, the correct procedure amounts to:
 1. Sampling a training and a testing set (from the same underlying distribution), e.g., with `make_gaussian_data(n, mu)`
 2. Fitting the classifier on training data
 3. Predicting the class labels of testing data
 4. Evaluating the fraction of wrong labels

```
x_tr, y_tr = make_gaussian_dataset(n, mu)
x_ts, y_ts = make_gaussian_dataset(n, mu)
clf = CNearestMeanClassifier()
clf.fit(x_tr, y_tr)
y_pred, dist = clf.predict(x_ts)
error = (y_pred != y_ts).mean()
```

What happens if one changes means and/or covariances of the Gaussian classes?
How does the error vary?

Lessons learned

- Visualize data and decision regions
- Implementation of a simple classifier (using a Python class)
- Create packages and dedicated function libraries
- Basic estimation of classifier performance on unseen data

Student challenges:

1. Extend `make_gaussian_dataset` to handle covariance matrices
2. Implement a k-Nearest Neighbor (kNN) classifier
3. Visualize decision regions of scikit-learn classifiers using
 - Nearest Centroid, and kNN (you may try other algorithms as well!)
4. Estimate performance of each classifier on unseen data

Please e-mail us if you are able to solve any of them!